

Q.1 Explain ACID properties of transaction

#answer

ACID Properties of Transactions:

1. **Atomicity:** The entire transaction is treated as a single unit. Either all operations in the transaction are completed or none are. If a failure occurs, the system rolls back to the state before the transaction started, ensuring no partial operations remain.
2. **Consistency:** A transaction takes the database from one consistent state to another. It ensures that any changes made by a transaction are valid according to predefined rules, constraints, and triggers.
3. **Isolation:** Transactions are executed in isolation from each other, ensuring that intermediate results of a transaction are not visible to other concurrent transactions. This prevents conflicts when multiple transactions are being processed simultaneously.
4. **Durability:** Once a transaction is committed, the changes it made are permanent, even in the event of a system crash. The system ensures that the results of committed transactions are not lost.

Q.2 Draw and explain state diagram of transactions

#answer

Q.3 Define Transaction with an example

#solvify :

- A transaction in a Database Management System (DBMS) is a sequence of operations performed as a single logical unit of work.
- A transaction ensures that the database moves from one consistent state to another even in the presence of system failures.
- Transactions are fundamental to maintaining the integrity and consistency of data in databases.

Example:

- In a banking system, a transaction might involve transferring money from one account to another. The transaction would include the following steps:
 - **Read** the balance of Account A.
 - **Deduct** the amount to be transferred from Account A.
 - **Add** the amount to Account B.
 - **Commit** the transaction, making the changes permanent.
- If any step fails (e.g., insufficient funds in Account A), the transaction would be rolled back to its initial state, ensuring no incorrect partial updates are made.

Q.4 Explain the steps involved in query processing and optimization.

#answer

Q.5 Give the advantages of concurrency in DBMS.

#answer

Q.6 Explain Timestamp Based Protocols

#answer

Q.7 Explain Validation based protocol.

#answer

Q.8 SQL query on subquery and join.

#answer

Q.9 Numerical on conflict serializability.

#answer

Q.10 Explain the concepts of LOCK in concurrency control.

#answer

Q.11 Explain starvation of transaction. Also state steps to avoid starvation.

#answer

- In DBMS (Database Management Systems), **starvation** refers to a situation where a particular transaction is perpetually delayed or never gets access to the resources it needs to proceed, even though the resources are available.
- This usually happens due to poor scheduling policies, leading to some transactions continuously being given lower priority

Example:

- Consider a scenario where multiple transactions are waiting to acquire a lock on a shared resource.
- If the DBMS continuously prioritizes one particular type of transaction (e.g., transactions with smaller data size), longer or lower-priority transactions may keep waiting indefinitely, even though the resource becomes available from time to time.

Solutions to Prevent Starvation:

❖ Fair Scheduling:

- Implementing fair scheduling algorithms such as **FIFO (First In, First Out)** or **Round Robin**, where each transaction gets a fair chance to execute, can help prevent starvation.

❖ Aging:

- The priority of a transaction can be dynamically increased as it waits in the queue.
- As time passes, the priority of a waiting transaction increases, ensuring that it eventually gets executed.

❖ Deadlock Detection and Resolution:

- Ensuring deadlock resolution strategies are fair, so that no transaction is constantly chosen for rollback, can also prevent starvation.

Q.12 Describe two-phase locking protocol.

#answer

- The **Two-Phase Locking Protocol (2PL)** is a concurrency control mechanism used in Database Management Systems (DBMS) to ensure that transactions are executed in a way that guarantees serializability.
- The protocol works by dividing the execution of a transaction into two distinct phases:
 1. **growing phase**
 2. **shrinking phase**

Growing Phase:

- In this phase, a transaction can acquire any number of locks (both shared and exclusive) on data items.
- However, it cannot release any locks.
- This phase continues until the transaction acquires all the locks it needs.

Shrinking Phase:

- Once the transaction releases its first lock, it enters the shrinking phase.
 - In this phase, the transaction can release locks but cannot acquire any new locks.
 - The transaction continues to release locks until it completes (either commits or aborts).
-

Q.13 Define deadlock with the help of an example.

#answer

Q.14 Explain wait-for graph for deadlock detection.

#answer

Q.15 Numerical on Timestamp Based Protocols

#answer

Q.16 Explain conflict serializability with example

#solvify :

- **Conflict Serializability:** Two schedules are conflict equivalent if they involve the same set of transactions and operations, and conflicting operations are in the same order.
- Conflicting operations are those that:
 - Belong to different transactions.
 - Access the same data item.
 - At least one operation is a write.
- If a schedule can be transformed into a **serial schedule** by swapping non-conflicting operations, it is **conflict serializable**.

Example:

T1	T2
Read(X)	
Write(X)	
	Read(X)
	Write(X)

- In this example:
 - **T1** reads the value of **X**, then writes a new value to **X**.
 - **T2** reads the updated value of **X** and then writes another value to **X**.
- Since both T1 and T2 are accessing the same data item (**X**) with at least one write operation, these operations conflict.

- The schedule can be transformed into a **serial schedule** where T1 completes all its operations before T2 starts. This schedule is conflict serializable because it is equivalent to a **serial execution**.

Q.17 Describe view serializability with example.

#solvify :

- **View Serializability:** Two schedules are view equivalent if:
 - The transactions in both schedules read the same initial values.
 - If a transaction in one schedule reads a value written by another transaction, the same holds true in the other schedule.
 - The final write operations in both schedules are the same.
- View serializability is a broader concept than conflict serializability. All conflict-serializable schedules are view serializable, but not all view-serializable schedules are conflict serializable.

Example of View Serializability:

T1	T2
Write(X)	
	Read(X)
	Write(Y)
Read(Y)	

- In this example, if T1 writes to **X** and T2 reads from **X** and writes to **Y**, the view of the final states in both schedules will maintain the same initial values and write operations, establishing view equivalence.

Q.18 Explain cascade and cascadeless schedule.

#answer

Q.19 Enlist the key features of NoSQL database

#answer

- **Schema Flexibility:** NoSQL databases often allow for dynamic schemas, meaning the structure of the data can change without needing to alter a predefined schema. This is particularly useful for applications where data types and structures are constantly evolving.
 - **Scalability:** Many NoSQL databases are designed to scale out horizontally by distributing data across multiple servers. This allows them to handle large volumes of data and high-velocity transactions more efficiently than vertical scaling (adding more power to a single server).
 - **High Performance:** NoSQL databases are optimized for performance with various data models and indexing techniques. They can handle high-speed reads and writes, making them suitable for real-time applications.
 - **High Availability and Fault Tolerance:** NoSQL databases are often designed with built-in replication and distribution features to ensure high availability and fault tolerance. This means they can continue to operate even if some parts of the system fail.
 - **Eventual Consistency:** Unlike relational databases that often use strong consistency models, many NoSQL databases use eventual consistency. This allows for high availability and partition tolerance, though it may mean that not all nodes have the most recent data at all times.
 - **Distributed Architecture:** Many NoSQL databases are built on a distributed architecture that allows data to be stored across multiple machines. This distributed nature helps manage large-scale data and provides resilience against failures.
 - **Easy Integration with Big Data Technologies:** NoSQL databases are frequently used alongside big data technologies such as Hadoop and Spark, providing a robust solution for handling large-scale data processing and analytics.
 - **Support for Unstructured Data:** They are well-suited for unstructured or semi-structured data, such as text, images, or social media posts, which may not fit neatly into traditional relational database tables.
 - **Flexible Query Languages:** NoSQL databases often come with their own query languages tailored to their specific data models, offering different querying capabilities compared to SQL.
-

Q.20 State advantages and disadvantages of using NoSQL databases.

#answer

Q.21 Differentiate between SQL and NoSQL databases briefly.

#answer

Sr. No	RDBMS	NoSQL
1	Relational or traditional database system	Non-relational or distributed database system
2	These databases have fixed or static schema	They have dynamic schema
3	These databases are not suited for hierarchical data storage	These databases are suited for hierarchical data storage
4	These databases are best suited for complex queries	These databases are not well suited for complex queries
5	Vertically scalable	Horizontally scalable
6	Follows ACID property	Follows CAP (consistency, availability, partition tolerance)
7	Examples: MySQL, Oracle, MS-SQL Server, etc	Examples: MongoDB, Cassandra, etc

Q.22 List two common types of NoSQL databases.

#answer

⚠ mentioned 4 types instead of 2

Types of NoSQL databases and the name of the database system that falls in that category are:

- **Graph Databases:** Examples – Amazon Neptune, Neo4j
- **Key value store:** Examples – Memcached, Redis, Coherence
- **Column:** Examples – Hbase, Big Table, Accumulo
- **Document-based:** Examples – MongoDB, CouchDB, Cloudant

Q.23 Give compatibility function.

#answer

- In the context of Database Management Systems (DBMS), a **compatibility function** is a mechanism used primarily in concurrency control to determine whether a set of operations (typically involving locks) can safely occur simultaneously without violating data integrity or consistency.
- The compatibility function in DBMS plays a crucial role in managing how transactions interact with each other through locks.
- By defining which lock types can coexist, it helps maintain data integrity, supports concurrency control, and aids in deadlock prevention.

Q.24 State the significance of SQL in Cassandra.

#answer

Q.25 Describe the basic steps involved in creating a new database in Cassandra.

#answer

Q.26 Describe the basic steps involved in creating a new database in MongoDB.

#answer

Q.27 Describe the steps involved in setting up a MongoDB database, creating collections, and inserting documents into a collection.

#answer
